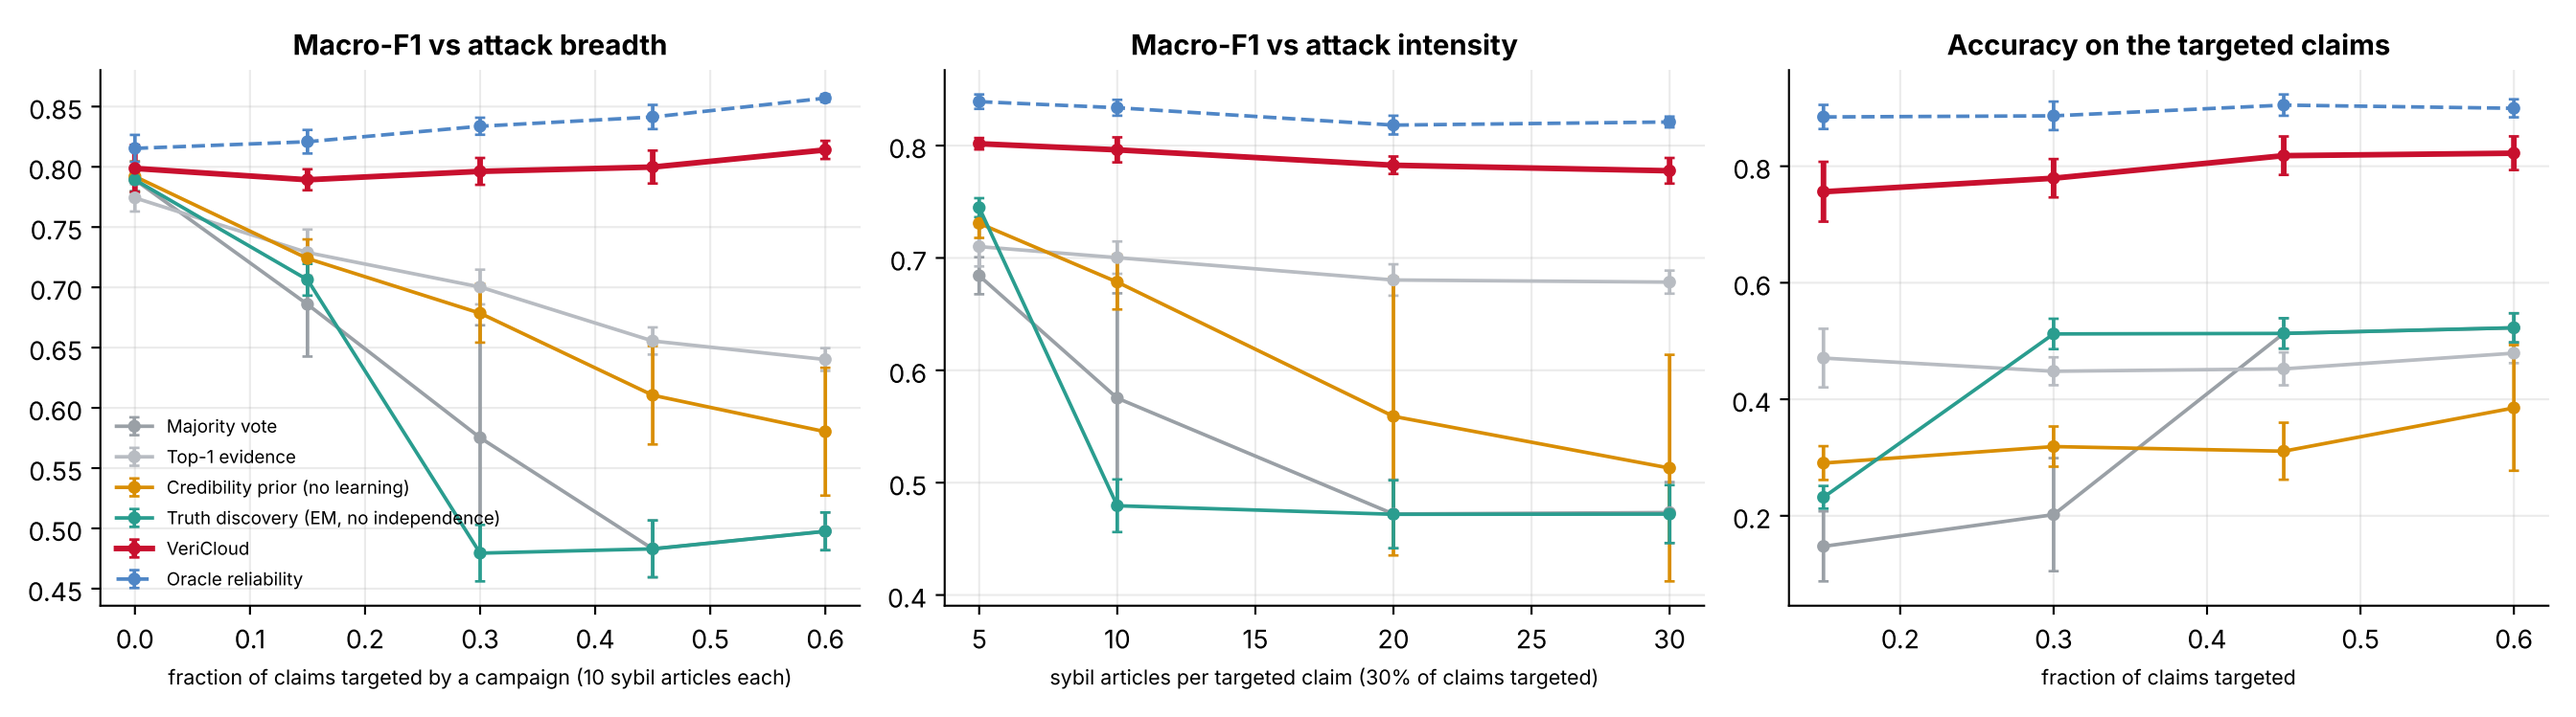
Macro-F1 vs attack breadth
0.85
0.80
0.75
0.70
0.65
0.60
0.55
0.50
0.45
Majority vote
Top-1 evidence
Credibility prior (no learning)
Truth discovery (EM, no independence)
VeriCloud
Oracle reliability
0.0
0.1
0.2
0.3
0.4
0.5
0.6
fraction of claims targeted by a campaign (10 sybil articles each)
Macro-F1 vs attack intensity
0.8
0.7
0.6
0.5
0.4
5
10
15
20
25
30
sybil articles per targeted claim (30% of claims targeted)
Accuracy on the targeted claims
0.8
0.6
0.4
0.2
0.2
0.3
0.4
0.5
0.6
fraction of claims targeted